

```
In [1]: 1 import numpy as np
        2 import pandas as pd
        3 import seaborn as sns
        4 import matplotlib.pyplot as plt
```

```
In [2]: 1 df=pd.read_csv('C:/Users/ADMIN/FIREBLAZE/FIREBLAZE ML/Black Friday Sales.csv')
        2 df.head()
```

Out[2]:

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status	Product_Category_1	Product_Catego
0	1000001	P00069042	F	0-17	10	A	2	0	3	
1	1000001	P00248942	F	0-17	10	A	2	0	1	
2	1000001	P00087842	F	0-17	10	A	2	0	12	
3	1000001	P00085442	F	0-17	10	A	2	0	12	
4	1000002	P00285442	M	55+	16	C	4+	0	8	



```
In [3]: 1 df.shape
```

Out[3]: (550068, 12)

In [4]: 1 df.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 550068 entries, 0 to 550067
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   User_ID                               550068 non-null  int64
1   Product_ID                            550068 non-null  object
2   Gender                                550068 non-null  object
3   Age                                    550068 non-null  object
4   Occupation                             550068 non-null  int64
5   City_Category                          550068 non-null  object
6   Stay_In_Current_City_Years            550068 non-null  object
7   Marital_Status                         550068 non-null  int64
8   Product_Category_1                     550068 non-null  int64
9   Product_Category_2                     376430 non-null  float64
10  Product_Category_3                     166821 non-null  float64
11  Purchase                               550068 non-null  int64
dtypes: float64(2), int64(5), object(5)
memory usage: 50.4+ MB

```

In [5]: 1 df.describe()

Out[5]:

	User_ID	Occupation	Marital_Status	Product_Category_1	Product_Category_2	Product_Category_3	Purchase
<b>count</b>	5.500680e+05	550068.000000	550068.000000	550068.000000	376430.000000	166821.000000	550068.000000
<b>mean</b>	1.003029e+06	8.076707	0.409653	5.404270	9.842329	12.668243	9263.968713
<b>std</b>	1.727592e+03	6.522660	0.491770	3.936211	5.086590	4.125338	5023.065394
<b>min</b>	1.0000001e+06	0.000000	0.000000	1.000000	2.000000	3.000000	12.000000
<b>25%</b>	1.001516e+06	2.000000	0.000000	1.000000	5.000000	9.000000	5823.000000
<b>50%</b>	1.003077e+06	7.000000	0.000000	5.000000	9.000000	14.000000	8047.000000
<b>75%</b>	1.004478e+06	14.000000	1.000000	8.000000	15.000000	16.000000	12054.000000
<b>max</b>	1.006040e+06	20.000000	1.000000	20.000000	18.000000	18.000000	23961.000000

## Checking Distribution

In [6]:

```
1 plt.figure(figsize=(7,4))
2 plt.title("Purchase Distribution")
3 sns.distplot(df['Purchase'],bins=25)
4
5 #First part of the graph has a normal distribution and later forming some peaks in the graph
6 #Evaluating the whole graph,it has a normal distribution
```

C:\Users\ADMIN\AppData\Local\Temp\ipykernel\_3368\370290696.py:3: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.


Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see


<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751> (<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>)

```
sns.distplot(df['Purchase'],bins=25)
```


Out[6]: <Axes: title={'center': 'Purchase Distribution'}, xlabel='Purchase', ylabel='Density'>




```
In [7]: 1 df['Gender'].value_counts()
2 plt.figure(figsize=(1,1))
3 sns.countplot(x='Gender', data=df, palette='husl')
4 plt.show()
5 #We observe that male customers have done more transactions than female during Black Friday sales.
```



```
In [8]: 1 plt.figure(figsize=(1,1))
2 sns.countplot(x='Marital_Status',data=df,palette='husl')
3 plt.show()
4 #There are more unmarried customers in the dataset who purchase more during Black Friday sales.
5 #Majority of the buyers are single
```



```
In [9]: 1 plt.figure(figsize=(1,1))
2 sns.countplot(x='City_Category',data=df,palette='husl')
3 plt.show()
4
5 #If we group by City_Category to see the variation with purchase
6 #We can see that city category B has made the most number of purchases followed by C.
7
8
```



In [10]:

```


1 plt.figure(figsize=(1,1))
2 df.groupby("City_Category").mean()["Purchase"].plot(kind='bar')
3 plt.title("City Category and Purchase Analysis")
4 plt.show()
5 #We then observe that the customers from city C spends the most

```

C:\Users\ADMIN\AppData\Local\Temp\ipykernel\_3368\4088996167.py:2: FutureWarning: The default value of numeric\_only in DataFrameGroupBy.mean is deprecated. In a future version, numeric\_only will default to False. Either specify numeric\_only or select only columns which should be valid for the function.

```
df.groupby("City_Category").mean()["Purchase"].plot(kind='bar')
```

City Category and Purchase Analysis




In [11]:


```

1 plt.figure(figsize=(1,1))
2 sns.countplot(x='Stay_In_Current_City_Years',data=df,palette='husl')
3 plt.show()
4 #It Looks Like the Longest customer is living in that city are less prone to buy new things
5 #whereas Customers which are new in town are more likely to take advantage of the low prices in Black Friday Sales

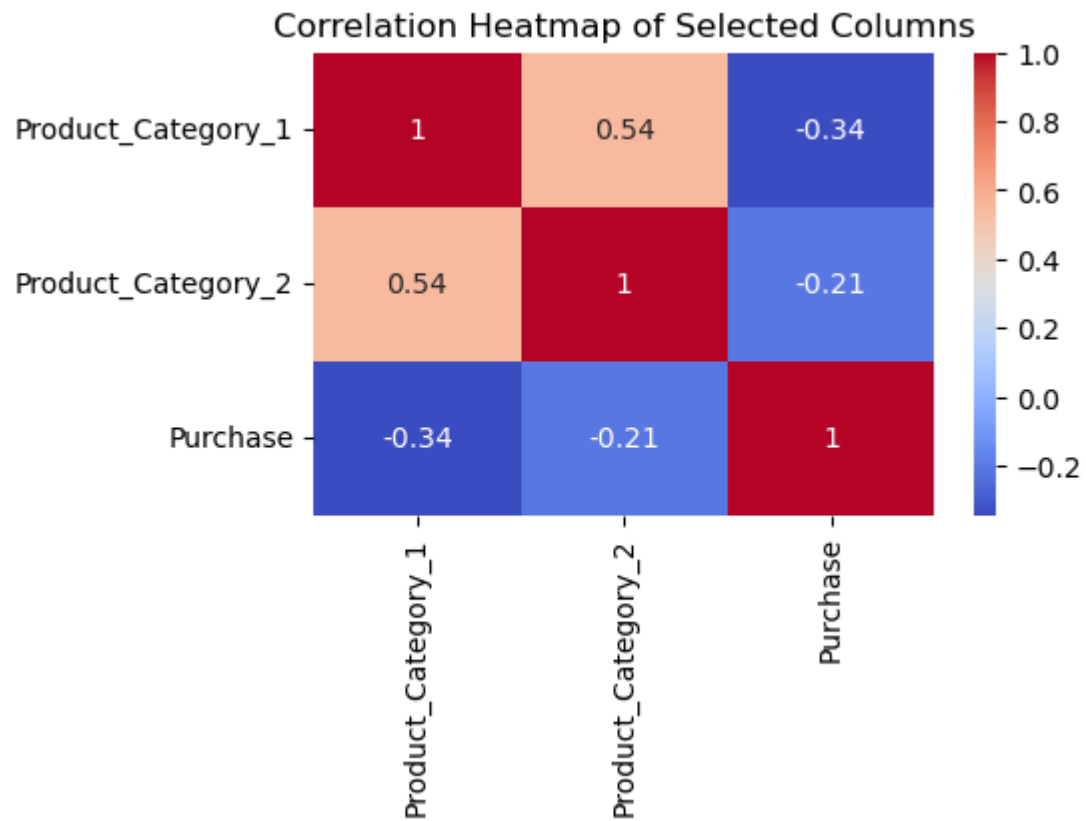
```



```
In [12]: 1 plt.figure(figsize=(4,4))
2 sns.countplot(x='Age',data=df,palette='husl')
3 plt.show()
4 #We can see that the Age group (26-35) makes the most no. of purchases during Black Friday Sales.
```



```
In [13]: 1 df1=df[["Product_Category_1","Product_Category_2","Purchase"]]
2 correlation=df1.corr()
3 plt.figure(figsize=(5,3))
4 sns.heatmap(correlation,annot=True,cmap="coolwarm")
5 plt.title("Correlation Heatmap of Selected Columns")
6 plt.show()
7
```





## Preprocessing dataset

```
In [14]: 1 df.isnull().sum()
```

```
Out[14]: User_ID          0
Product_ID          0
Gender              0
Age                0
Occupation          0
City_Category      0
Stay_In_Current_City_Years  0
Marital_Status      0
Product_Category_1  0
Product_Category_2  173638
Product_Category_3  383247
Purchase            0
dtype: int64
```

```
In [15]: 1 Product_Category_3=(df["Product_Category_3"].isnull().sum()/len(df["Product_Category_3"]))*100
2 Product_Category_3
3 #the mentioned column has null values more than 50%
```

```
Out[15]: 69.67265865311198
```

```
In [16]: 1 df=df.drop(columns=["Product_Category_3"])
```

```
In [17]: 1 df['Product_Category_2'].fillna(df['Product_Category_2'].mean(),inplace=True)
```

```
In [18]: 1 df.nunique()
```

```
Out[18]: User_ID          5891
Product_ID        3631
Gender              2
Age                7
Occupation         21
City_Category      3
Stay_In_Current_City_Years  5
Marital_Status     2
Product_Category_1  20
Product_Category_2  18
Purchase          18105
dtype: int64
```

```
In [19]: 1 df=df.drop(["User_ID","Product_ID"],axis=1)
2
3 #We can drop User_ID and Product_ID for model prediction as it has more unique values Or
4 #if not the results will be biased to User_ID or Product_ID
5
```

## Converting categorical data into numerical

```
In [20]: 1 from sklearn.preprocessing import LabelEncoder
```


```
In [21]: 1 LE=LabelEncoder()
```


```
In [22]: 1 categorical_columns=['Gender','Age','City_Category','Stay_In_Current_City_Years']
```


```
In [23]: 1 for column in categorical_columns:
2         df[column] = LE.fit_transform(df[column])
3
```

## Checking skewness

```
In [24]: 1 for i in df1:
2         sns.kdeplot(data=df1,x=i)
3         plt.show()
4
5         #from this we can see Product_Category_1 is rightly skewed
```







## Treating skewness of column with PowerTransformer

```
In [25]: 1 from sklearn.preprocessing import PowerTransformer
```


```
In [26]: 1 PT=PowerTransformer()
```


```
In [27]: 1 columns_to_scale=['Product_Category_1','Product_Category_2','Purchase']
2 for column in columns_to_scale:
3     df[column] = PT.fit_transform(df[[column]])
```

## Checking & Treating outliers


```
In [28]: 1 lst=df[['Product_Category_1','Purchase']]
2 for col in lst:
3     Q3=df[col].quantile(0.75)
4     Q1=df[col].quantile(0.25)
5     IQR=Q3-Q1
6     uw=Q3+1.5*IQR
7     lw=Q1-1.5*IQR
8     df=df[(df[col]>=lw)&(df[col]<=uw)]
```

```
In [29]: 1 for i in df1:  
2         sns.boxplot(data=df,x=i)  
3         plt.show()
```









## Split the data for training & testing

```
In [30]: 1 from sklearn.model_selection import train_test_split
```

```
In [31]: 1 x=df.drop('Purchase',axis=1)
         2 y=df['Purchase']
```

```
In [32]: 1 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.30,random_state=0)
```

## LinearRegression

```
In [33]: 1 from sklearn.linear_model import LinearRegression,Lasso
```

```
In [34]: 1 LR=LinearRegression().fit(x_train,y_train)
```

```
In [35]: 1 LR.score(x_train,y_train)
```

```
Out[35]: 0.18770518525090651
```

```
In [36]: 1 y_pred=LR.predict(x_test)
```

```
In [37]: 1 from sklearn.metrics import r2_score,mean_squared_error,mean_absolute_error
```

```
In [38]: 1 mse=mean_squared_error(y_test,y_pred)
2 mae=mean_absolute_error(y_test,y_pred)
3 r2_score(y_test,y_pred)
4
5 print("r2_score:          ",(r2_score(y_test,y_pred)))
6 print("mean_squared_error: ",(mean_squared_error(y_test,y_pred)))
7 print("mean_absolute_error: ",(mean_absolute_error(y_test,y_pred)))
```

```
r2_score:          0.1815723738838485
mean_squared_error: 0.7774466714826213
mean_absolute_error: 0.6637134352307733
```

```
In [39]: 1 residual=y_test-y_pred #Residuals are the differences between the observed values and the values predicted by model
```

## DecisionTreeRegressor

```
In [40]: 1 from sklearn.tree import DecisionTreeRegressor
```

```
In [41]: 1 DTR=DecisionTreeRegressor(random_state=0).fit(x_train,y_train)
```

```
In [42]: 1 DTR.score(x_train,y_train)
```

```
Out[42]: 0.7484293563655522
```

```
In [43]: 1 y_pred=DTR.predict(x_test)
```

```
In [44]: 1 r2_score(y_test,y_pred)
```

```
Out[44]: 0.6053220765476833
```

```
In [45]: 1 print("r2_score:          ",(r2_score(y_test,y_pred)))  
2 print("mean_squared_error:  ",(mean_squared_error(y_test,y_pred)))  
3 print("mean_absolute_error: ",(mean_absolute_error(y_test,y_pred)))
```

```
r2_score:          0.6053220765476833  
mean_squared_error: 0.3749152986829032  
mean_absolute_error: 0.45428613008948127
```

## RandomForestRegressor

```
In [46]: 1 from sklearn.ensemble import RandomForestRegressor
```

```
In [47]: 1 RFR=RandomForestRegressor(n_estimators=100,max_depth=15,random_state=0).fit(x_train,y_train)
```

```
In [48]: 1 DTR.score(x_train,y_train)
```

```
Out[48]: 0.7484293563655522
```

```
In [49]: 1 y_pred=RFR.predict(x_test)
```

```
In [50]: 1 print("r2_score:          ",(r2_score(y_test,y_pred)))
2 print("mean_squared_error:  ",(mean_squared_error(y_test,y_pred)))
3 print("mean_absolute_error: ",(mean_absolute_error(y_test,y_pred)))
```

```
r2_score:          0.6593111530508973
mean_squared_error: 0.3236296058684633
mean_absolute_error: 0.4357836320832321
```

## KNeighborsRegressor

```
In [51]: 1 from sklearn.neighbors import KNeighborsRegressor
```

```
In [52]: 1 KNR=KNeighborsRegressor(n_neighbors=15).fit(x_train,y_train)
```

```
In [53]: 1 KNR.score(x_train,y_train)
```

```
Out[53]: 0.48710993328704766
```

```
In [54]: 1 y_pred=KNR.predict(x_test)
```

```
In [55]: 1 print("r2_score:          ",(r2_score(y_test,y_pred)))
2 print("mean_squared_error:  ",(mean_squared_error(y_test,y_pred)))
3 print("mean_absolute_error: ",(mean_absolute_error(y_test,y_pred)))
```

```
r2_score:          0.41852151822054695
mean_squared_error: 0.5523622318854204
mean_absolute_error: 0.5494539115276138
```

In [ ]:

1